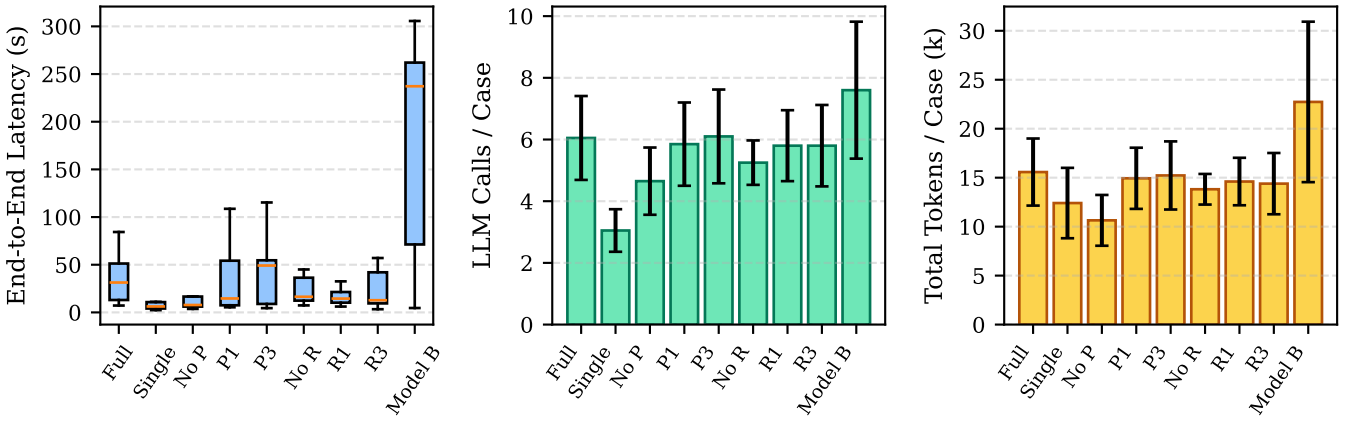
(a) Per-Case Latency (s)
End-to-End Latency (s)
0
50
100
150
200
250
300
Full
Single
No P
P1
P3
No R
R1
R3
Model B
(b) Mean LLM Calls
LLM Calls / Case
0
2
4
6
8
10
Full
Single
No P
P1
P3
No R
R1
R3
Model B
(c) Mean Total Tokens (k)
Total Tokens / Case (k)
0
5
10
15
20
25
30
Full
Single
No P
P1
P3
No R
R1
R3
Model B