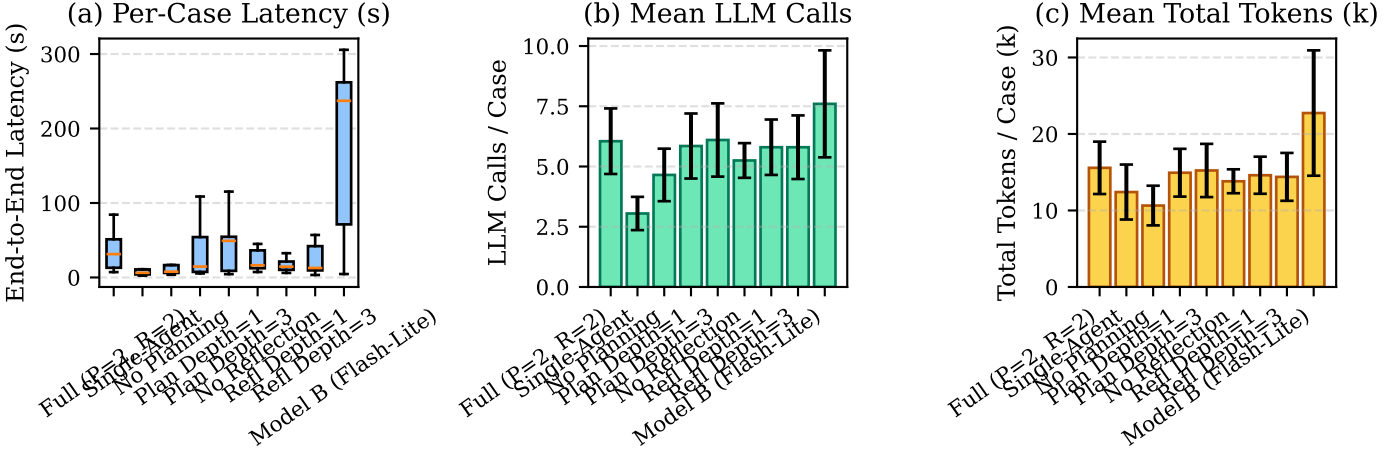
(a) Per-Case Latency (s)
End-to-End Latency (s)
0
100
200
300
(b) Mean LLM Calls
LLM Calls / Case
0.0
2.5
5.0
7.5
10.0
(c) Mean Total Tokens (k)
Total Tokens / Case (k)
0
10
20
30
Full (P=2, R=2)
Single-Agent
No Planning
Plan Depth=1
Plan Depth=3
No Reflection
Refl Depth=1
Refl Depth=3
Model B (Flash-Lite)