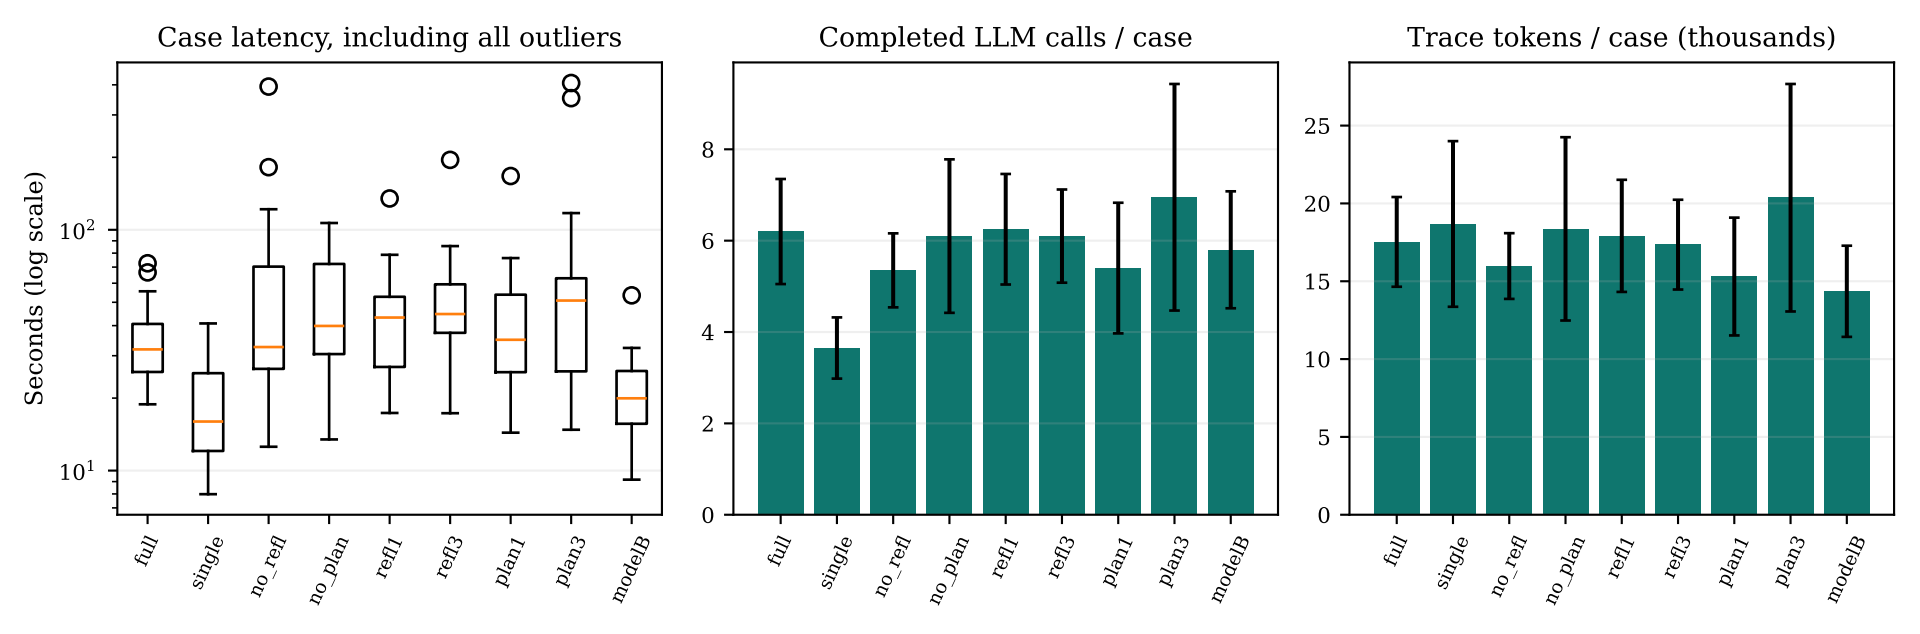
Case latency, including all outliers
Seconds (log scale)
10^2
10^1
full
single
no_refl
no_plan
refl1
refl3
plan1
plan3
modelB
Completed LLM calls / case
0
2
4
6
8
Trace tokens / case (thousands)
0
5
10
15
20
25